

[🏠](#) / [Examples](#) / [Orchestration Service](#)

Orchestration Service ¶

This notebook demonstrates how to use the SDK to interact with the Orchestration Service, enabling the creation of AI-driven workflows by seamlessly integrating various modules, such as templating, large language models (LLMs), data masking and content filtering. By leveraging these modules, you can build complex, automated workflows that enhance the capabilities of your AI solutions. For more details on configuring and using these modules, please refer to the [Orchestration Service Documentation](#).

Prerequisite ¶

! Important: Before you begin using the SDK, make sure to set up a virtual deployment of the Orchestration Service. This deployment process will provide you with a unique endpoint URL (**deploymentUrl**) for your instance.

For detailed guidance on setting up the Orchestration Service, please refer to the setup guide [here](#).

```
YOUR_API_URL = "... " # deployment URL of the orchestration service
```

Basic Orchestration Pipeline ¶

Now that you have `YOUR_API_URL`, let's walk through a basic orchestration pipeline for a translation task.

Step 1: Define the Template and Default Input Values ¶

The `Template` class is used to define structured message templates for generating dynamic interactions with language models. In this example, the template is designed for a

```
from gen_ai_hub.orchestration.models.message import SystemMessage, UserMessage
from gen_ai_hub.orchestration.models.template import Template, TemplateValue

template = Template(
    messages=[
        SystemMessage("You are a helpful translation assistant."),
        UserMessage(
            "Translate the following text to {{to_lang}}: {{text}}"
        ),
    ],
    defaults=[
        TemplateValue(name="to_lang", value="German"),
    ],
)
```

This template can be used to create translation requests where the language and text to be translated are specified dynamically. The placeholders in the `UserMessage` will be replaced with the actual values provided at runtime, and the default value for the language is set to German.

Referencing templates in the prompt registry ¶

You can also use a prompt template reference. This approach allows you to reuse existing templates stored in the Prompt Registry.

```
from gen_ai_hub.orchestration.models.template_ref import TemplateRef

template_by_id = TemplateRef.from_id(prompt_template_id="648871d9-b207-4
template_by_names = TemplateRef.from_tuple(scenario="translation", name=
```

Step 2: Define the LLM ¶

The `LLM` class is used to configure and initialize a language model for generating text based on specific parameters. In this example, we'll use the `gpt-4o` model to perform the translation task.

Note: The Orchestration Service automatically manages the virtual deployment of the language model, so no additional setup is needed on your end.

```
from gen_ai_hub.orchestration.models.llm import LLM
```

This configuration initializes the language model to use the `gpt-4o` model with the latest version. The model will generate responses up to 256 tokens in length and produce more predictable and focused output due to the low temperature setting.

Step 3: Create the Orchestration Configuration ¶

The `OrchestrationConfig` class defines a configuration for integrating various modules, such as templates and language models, into a cohesive orchestration setup. It specifies how these components interact and are configured to achieve the desired operational scenario.

```
from gen_ai_hub.orchestration.models.config import OrchestrationConfig

config = OrchestrationConfig(
    template=template, # or use a referenced prompt template from Step
    llm=llm,
)
```

Step 4: Run the Orchestration Request ¶

The `OrchestrationService` class is used to interact with a specific orchestration service instance by providing configuration details, including the `API URL`, to initiate and manage its operations.

```
from gen_ai_hub.orchestration.service import OrchestrationService

orchestration_service = OrchestrationService(api_url=YOUR_API_URL, confi
```

Call the `run` method with the required `template values`. The service will process the input according to the configuration and return the result.

```
result = orchestration_service.run(template_values=[
    TemplateValue(name="text", value="The Orchestration Service is worki
])
print(result.orchestration_result.choices[0].message.content)
```

Der Orchestrierungsdienst funktioniert!

Optional Modules ¶

Data Masking ¶

The Data Masking module anonymizes or pseudonymizes personally identifiable information (PII) before it is processed by the LLM module. When data is anonymized, all identifying information is replaced with placeholders (e.g., MASKED_ENTITY), and the original data cannot be recovered, ensuring that no trace of the original information is retained. In contrast, pseudonymized data is substituted with unique placeholders (e.g., MASKED_ENTITY_ID), allowing the original information to be restored if needed. In both cases, the masking module identifies sensitive data and replaces it with appropriate placeholders before further processing.

```
from gen_ai_hub.orchestration.utils import load_text_file
from gen_ai_hub.orchestration.models.data_masking import DataMasking
from gen_ai_hub.orchestration.models.sap_data_privacy_integration import ProfileEntity

data_masking = DataMasking(
    providers=[
        SAPDataPrivacyIntegration(
            method=MaskingMethod.ANONYMIZATION, # or MaskingMethod.PSEUDONYMIZATION
            entities=[
                ProfileEntity.EMAIL,
                ProfileEntity.PHONE,
                ProfileEntity.PERSON,
                ProfileEntity.ORG,
                ProfileEntity.LOCATION
            ]
        )
    ]
)

config = OrchestrationConfig(
    template=Template(
        messages=[
            SystemMessage("You are a helpful AI assistant."),
            UserMessage("Summarize the following CV in 10 sentences: {{?
        ]
    ),
    llm=LLM(
        name="gpt-4o",
    ),
    data_masking=data_masking
)
```

```
cv_as_string = load_text_file("data/cv.txt")

result = orchestration_service.run(
    config=config,
    template_values=[
        TemplateValue(name="orgCV", value=cv_as_string)
    ]
)
```

```
print(result.orchestration_result.choices[0].message.content)
```

The CV belongs to a financial professional with a strong background in

Content Filtering ¶

The Content Filtering module can be configured to filter both the input to the LLM module (input filter) and the output generated by the LLM (output filter). The module uses predefined classification services to detect inappropriate or unwanted content, allowing flexible configuration through customizable thresholds. These thresholds can be set to control the sensitivity of filtering, ensuring that content meets desired standards before it is processed or returned as output.

```
from gen_ai_hub.orchestration.models.content_filtering import InputFilter
from gen_ai_hub.orchestration.models.azure_content_filter import AzureContentFilter

input_filter= AzureContentFilter(hate=AzureThreshold.ALLOW_SAFE,
                                violence=AzureThreshold.ALLOW_SAFE,
                                self_harm=AzureThreshold.ALLOW_SAFE,
                                sexual=AzureThreshold.ALLOW_SAFE)
output_filter = AzureContentFilter(hate=AzureThreshold.ALLOW_SAFE,
                                   violence=AzureThreshold.ALLOW_SAFE_LO,
                                   self_harm=AzureThreshold.ALLOW_SAFE_LO,
                                   sexual=AzureThreshold.ALLOW_ALL)

config = OrchestrationConfig(
    template=Template(
        messages=[
            SystemMessage("You are a helpful AI assistant."),
            UserMessage("{{?text}}"),
        ]
    ),
    llm=LLM(
        name="gpt-4o",
```

```
input_filtering=InputFiltering(filters=[input_filter]),
output_filtering=OutputFiltering(filters=[output_filter])
)
```

```
from gen_ai_hub.orchestration.exceptions import OrchestrationError

try:
    result = orchestration_service.run(config=config, template_values=[
        TemplateValue(name="text", value="I hate you")
    ])
except OrchestrationError as error:
    print(error.message)
```

Content filtered due to safety violations. Please modify the prompt and

Streaming ¶

When you initiate an orchestration request, the full response is typically processed and delivered in one go. For longer responses, this can lead to delays in receiving the complete output. To mitigate this, you have the option to stream the results as they are being generated. This helps in rapidly processing or displaying initial portions of the results without waiting for the entire computation to finish.

To activate streaming, use the `stream` method of the `OrchestrationService`. This method returns an object that streams chunks of the response as they become available. You can then extract relevant information from the `delta` field.

Here's how you can set up a simple configuration to stream orchestration results:

```
config = OrchestrationConfig(
    template=Template(
        messages=[
            SystemMessage("You are a helpful AI assistant."),
            UserMessage("{{?text}}"),
        ]
    ),
    llm=LLM(
        name="gpt-35-turbo",
        parameters={
            "max_tokens": 256,
```

```

        }
    ),
)

service = OrchestrationService(
    api_url=YOUR_API_URL
)

response = service.stream(
    config=config,
    template_values=[
        TemplateValue(name="text", value="Which color is the sky? Answer")
    ]
)

for chunk in response:
    print(chunk.orchestration_result)
    print("*" * 20)

```

```

LLMResultStreaming(id='', object='', created=0, model='', choices=[LLM
*****
LLMResultStreaming(id='chatcmpl-AZyDceGqjGjvHxAtDqTZZgKn2azbn', object
*****
LLMResultStreaming(id='chatcmpl-AZyDceGqjGjvHxAtDqTZZgKn2azbn', object
*****

```

Note: As shown above, streaming responses contain a delta field instead of a message field.

You can customize the global stream behavior by setting options like `chunk_size` which controls the amount of data processed in each chunk:

```

response = service.stream(
    config=config,
    template_values=[
        TemplateValue(name="text", value="Which color is the sky? Answer")
    ],
    stream_options={
        'chunk_size': 25
    }
)

for chunk in response:
    print(chunk.orchestration_result)
    print("*" * 20)

```

```

LLMResultStreaming(id='', object='', created=0, model='', choices=[LLM
*****
LLMResultStreaming(id='chatcmpl-AZyDduknxgCQtpmQAvgPyu7mgCS9n', object
*****
LLMResultStreaming(id='chatcmpl-AZyDduknxgCQtpmQAvgPyu7mgCS9n', object
*****
LLMResultStreaming(id='chatcmpl-AZyDduknxgCQtpmQAvgPyu7mgCS9n', object
*****
LLMResultStreaming(id='chatcmpl-AZyDduknxgCQtpmQAvgPyu7mgCS9n', object
*****
LLMResultStreaming(id='chatcmpl-AZyDduknxgCQtpmQAvgPyu7mgCS9n', object
*****
LLMResultStreaming(id='chatcmpl-AZyDduknxgCQtpmQAvgPyu7mgCS9n', object
*****

```

Modules that influence or process streaming results, such as `OutputFiltering`, might need specific stream options. The `overlap` option allows you to include extra context during the filtering process:

```

config = OrchestrationConfig(
    template=Template(
        messages=[
            SystemMessage("You are a helpful AI assistant."),
            UserMessage("{?text}"),
        ]
    ),
    llm=LLM(
        name="gpt-35-turbo",
        parameters={
            "max_tokens": 256,
            "temperature": 0.0
        }
    ),
    output_filtering=OutputFiltering(
        filters=[
            AzureContentFilter(
                hate=AzureThreshold.ALLOW_ALL,
                violence=AzureThreshold.ALLOW_ALL,
                self_harm=AzureThreshold.ALLOW_ALL,
                sexual=AzureThreshold.ALLOW_ALL
            )
        ],
        stream_options={
            'overlap': 100
        }
    )
)

```



```

        config=config,
        template_values=[
            TemplateValue(name="text", value="Why is the sky blue?")
        ]
    )

    for chunk in response:
        print(chunk.orchestration_result.choices[0].delta.content, end='')

```

The blue color of the sky is due to a phenomenon called Rayleigh scatt

Advanced Examples ¶

```
service = OrchestrationService(api_url=YOUR_API_URL)
```

Translation Service ¶

This example extends the initial walkthrough of a basic orchestration pipeline by abstracting the translation task into its own reusable `TranslationService` class. Once the configuration is established, it can be easily adapted and reused for different translation scenarios.

```

from gen_ai_hub.orchestration.models.config import OrchestrationConfig
from gen_ai_hub.orchestration.models.llm import LLM
from gen_ai_hub.orchestration.models.message import SystemMessage, UserMessage
from gen_ai_hub.orchestration.models.template import Template, TemplateValue
from gen_ai_hub.orchestration.service import OrchestrationService

class TranslationService:
    def __init__(self, orchestration_service: OrchestrationService):
        self.service = orchestration_service
        self.config = OrchestrationConfig(
            template=Template(
                messages=[
                    SystemMessage("You are a helpful translation assistant"),
                    UserMessage(
                        "Translate the following text to {{to_lang}}: {{text}}"
                    ),
                ],
                defaults=[
                    TemplateValue(name="to_lang", value="English"),

```

```
        ),
        llm=LLM(name="gpt-4o"),
    )

    def translate(self, text, to_lang):
        response = self.service.run(
            config=self.config,
            template_values=[
                TemplateValue(name="to_lang", value=to_lang),
                TemplateValue(name="text", value=text),
            ],
        )

        return response.orchestration_result.choices[0].message.content
```

```
translator = TranslationService(orchestration_service=service)
```

```
result = translator.translate(text="Hello, world!", to_lang="French")
print(result)
```

Bonjour, le monde !

```
result = translator.translate(text="Hello, world!", to_lang="Spanish")
print(result)
```

¡Hola, mundo!

```
result = translator.translate(text="Hello, world!", to_lang="German")
print(result)
```

Hallo, Welt!

Chatbot with Memory ¶

This example demonstrates how to integrate the `OrchestrationService` with a chatbot to handle conversational flow.

When making requests to the orchestration service, you can specify a list of messages as `history` that will be prepended to the templated content and processed by the templating module. These messages are plain, non-templated messages, as they typically represent past conversation outputs — such as in this chatbot scenario.



Search ...



```
from typing import List

from gen_ai_hub.orchestration.models.config import OrchestrationConfig
from gen_ai_hub.orchestration.models.llm import LLM
from gen_ai_hub.orchestration.models.message import Message, SystemMessage, UserMessage
from gen_ai_hub.orchestration.models.template import Template, TemplateValue
from gen_ai_hub.orchestration.service import OrchestrationService


class ChatBot:
    def __init__(self, orchestration_service: OrchestrationService):
        self.service = orchestration_service
        self.config = OrchestrationConfig(
            template=Template(
                messages=[
                    SystemMessage("You are a helpful chatbot assistant."),
                    UserMessage("{{?user_query}}"),
                ],
            ),
            llm=LLM(name="gpt-4"),
        )
        self.history: List[Message] = []

    def chat(self, user_input):
        response = self.service.run(
            config=self.config,
            template_values=[
                TemplateValue(name="user_query", value=user_input),
            ],
            history=self.history,
        )

        message = response.orchestration_result.choices[0].message

        self.history = response.module_results.templating
        self.history.append(message)

        return message.content

    def reset(self):
        self.history = []
```

```
bot = ChatBot(orchestration_service=service)
```

```
print(bot.chat("Hello, how are you?"))
```

Hello! I'm just a chatbot, so I don't have feelings, but I'm here to h

```
print(bot.chat("What's the weather like today?"))
```

I don't have real-time data capabilities. I recommend checking a relia

```
print(bot.chat("Can you remember what I first asked you?"))
```

Yes, you first asked me how I am doing. How can I further assist you t

```
bot.reset()
```

```
print(bot.chat("Can you remember what I first asked you?"))
```

As an AI model developed by OpenAI, I don't have memory in our current

Sentiment Analysis with Few Shot Learning ¶

This example demonstrates the different message roles in the templating module through a few-shot learning use case with the `FewShotLearner` class.

- **Message Types:** Different message types (`SystemMessage` , `UserMessage` , `AssistantMessage`) structure the interaction and guide the model's behavior.
- **Templating:** The template includes these examples, ending with a placeholder (`{{? user_input}}`) for dynamic user input.
- **Few-Shot Examples:** Pairs of `UserMessage` and `AssistantMessage` show how the

The FewShotLearner class manages the dynamic creation of the template and ensures the correct message roles are used for each user input.

```

from typing import List, Tuple

from gen_ai_hub.orchestration.models.config import OrchestrationConfig
from gen_ai_hub.orchestration.models.llm import LLM
from gen_ai_hub.orchestration.models.message import (
    SystemMessage,
    UserMessage,
    AssistantMessage,
)
from gen_ai_hub.orchestration.models.template import Template, TemplateValue
from gen_ai_hub.orchestration.service import OrchestrationService


class FewShotLearner:
    def __init__(
        self,
        orchestration_service: OrchestrationService,
        system_message: SystemMessage,
        examples: List[Tuple[UserMessage, AssistantMessage]],
    ):
        self.service = orchestration_service
        self.config = OrchestrationConfig(
            template=self._create_few_shot_template(system_message, exam
            llm=LLM(name="gpt-35-turbo"),
        )

    @staticmethod
    def _create_few_shot_template(
        system_message: SystemMessage,
        examples: List[Tuple[UserMessage, AssistantMessage]],
    ) → Template:
        messages = [system_message]

        for example in examples:
            messages.append(example[0])
            messages.append(example[1])
        messages.append(UserMessage("{?user_input}"))

        return Template(messages=messages)

    def predict(self, user_input: str) → str:
        response = self.service.run(
            config=self.config,
            template_values=[TemplateValue(name="user_input", value=user
        )

        return response.orchestration_result.choices[0].message.content

```

```
sentiment_examples = [  
    (UserMessage("I love this product!"), AssistantMessage("Positive")),  
    (UserMessage("This is terrible service."), AssistantMessage("Negative")),  
    (UserMessage("The weather is okay today."), AssistantMessage("Neutral"))  
]
```

```
sentiment_analyzer = FewShotLearner(  
    orchestration_service=service,  
    system_message=SystemMessage(  
        "You are a sentiment analysis assistant. Classify the sentiment of the following text."  
    ),  
    examples=sentiment_examples,  
)
```

```
print(sentiment_analyzer.predict("The movie was a complete waste of time"))
```

Negative

```
print(  
    sentiment_analyzer.predict("The traffic was fortunately unusually light today.")  
)
```

Positive

```
print(  
    sentiment_analyzer.predict("I'm not sure how I feel about the recent election results.")  
)
```

Neutral